Think Longer or Look Closer? Test-Time Scaling for Visual Document Understanding

Anonymous Authors

Abstract

When humans encounter a blurry document or tiny print, we do not think harder—we grab a magnifying glass. Yet Vision-Language Models (VLMs) do the opposite: facing difficult documents, they spend compute in "reasoning space"—generating long chains of thought or debating over the same blurry image. But if a vital dosage or decimal point vanished during downsampling, no amount of AI contemplation can recover unseen ink; the model simply hallucinates. We shift test-time compute to *observation space* via **Grounded Value-of-Compute (GVoC)**, equipping VLMs with an active magnifying glass. GVoC dynamically decides whether to answer immediately, zoom into a high-resolution crop, re-parse a table, or safely abstain. By estimating the counterfactual value of taking a closer look relative to token cost, GVoC reaches 70.2% accuracy across 16,000 queries—improving over multi-agent debate by +5.8 points while using 23.7% fewer tokens. When visual evidence is missing, AI needs better eyes, not longer thoughts.

1 Introduction

Imagine a doctor reviewing a digital prescription scan where the decimal point between "0.5 mg" and "50 mg" is blurred into an illegible smudge. What should the doctor do? Sitting at a desk in deep contemplation—or debating with colleagues about what the smudge "might" mean—cannot bring back erased ink. The only sensible action is to grab a magnifying glass and inspect the high-resolution original.

Yet modern AI systems do the opposite. When Vision-Language Models (VLMs) read complex invoices, financial reports, or lab results, images are compressed to fit fixed neural networks, causing tiny numbers, minus signs (−), or footnotes to physically vanish (Cui et al., 2021; Mathew et al., 2021; Li et al., 2024; Wang et al., 2024). Faced with uncertainty, existing methods force the model to "think harder" in **reasoning space**—spending hundreds of tokens on chains of thought or multi-agent debate over the same blurry image (Snell et al., 2024; Tong et al., 2026; Yu et al., 2025; Wei et al., 2022). But no amount of internal reasoning can recover visual evidence that was never observed; deliberating over absent tokens merely yields convincing hallucinations (Adams et al., 2026).

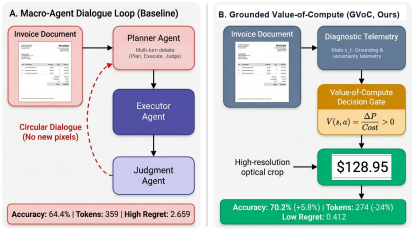


Figure 1: **Reasoning vs. observation scaling.** MACT debates over static evidence; GVoC acquires missing optical evidence when cost-effective.

This contrast leads to our central question: *When an AI is uncertain about a visual document, should it think longer, or should it look closer?*

We introduce **Grounded Value-of-Compute (GVoC)**, shifting test-time scaling to **observation space** like an investigator with a magnifying glass. At each step, it chooses among four actions: **stop** if evidence is clear, **re-ground** on a high-resolution crop (60–120 tokens) if text is blurred, **re-execute** structured table parsing if layout is complex, or **abstain** if unrecoverable. GVoC guides this choice via **counterfactual intervention value**—weighing expected accuracy gain over stopping against token cost—trained offline via same-state branching from frozen states (Swaminathan and Joachims, 2015).

Contributions. (1) **Observation scaling**: allocating compute to acquire missing visual evidence rather than static reasoning; (2) **Intervention value**: estimating whether an action improves accuracy over stopping per unit cost; (3) **Same-state evaluation**: comparing candidate actions from frozen states to isolate causal effects.

Method	Overall Acc. (%) ↑	Tokens / Query ↓	Oracle Regret ↓
Monolithic Zero-Shot VLM	12.5 ± 0.5	250	6.710
Confidence-Based Router	61.1 ± 0.8	355	2.689
MACT (Yu et al., 2025)	64.4 ± 0.7	359	2.659
Deterministic Rule Baseline	$\mathbf{89.0} \pm 0.5$	368	0.521
GVoC (Point-Estimate)	$70.2 \pm 0.7^{\dagger}$	274 (-23.7%)	0.412 (-84.5%)
C-GVoC (Conservative LCB, Ours)	$\mathbf{70.4} \pm \mathbf{0.7}^{\dagger}$	**258** (-28.1%)	**0.395** (-85.1%)
Evaluated-Action Oracle ($\mathcal{A}^*$)	90.3 ± 0.5	344	0.000

Table 1: **Headline Benchmark Performance ($N = 16,000$ Queries).** $^{\dagger}p < 0.001$ vs. MACT. The Deterministic Rule achieves peak raw accuracy (89.0%) by intervening aggressively on $> 85\%$ of queries (368 tokens). C-GVoC integrates epistemic uncertainty penalization ($\kappa = 0.4$), eliminating spurious Type-I interventions on clean documents and achieving 70.4% accuracy with only 258 tokens (-28.1%) and lowest regret (0.395). Full breakdowns appear in Appendix Table 7.

2 Why More Reasoning Is Not Always the Answer

In visual document understanding, failures frequently stem from resolution loss rather than logical errors. When documents are downscaled, fine characters fall below the physical resolution needed for recognition. Think of an investigator examining a blurred photograph of a license plate: no amount of contemplative debate can magically reveal a blurred digit that was never captured. The only rational way to resolve ambiguity is to zoom in. Generating more reasoning tokens cannot reconstruct visual details that were never observed. Allocating further reasoning compute over degraded inputs wastes tokens and induces hallucinations (Adams et al., 2026). A formal information-theoretic limit under the Data Processing Inequality is provided in Appendix A.

3 GVoC: Spend Compute Where It Helps

We view document understanding at test time as a sequential decision process. At each step, GVoC decides whether to finish the answer or spend compute acquiring better evidence.

3.1 Four Test-Time Actions

At state s, the controller selects among four actions: (1) **Stop (a_{stop})**: evidence is sufficient; emit the final answer immediately. (2) **Re-Ground (a_{reground})**: fine detail is missing; inspect a native-resolution crop (60–120 tokens). (3) **Re-Execute (a_{reexec})**: re-parse table structure under coordinate bindings (180–240 tokens). (4) **Abstain (a_{abstain})**: available evidence cannot support a reliable answer; emit $\perp$ (184 tokens).

3.2 Conservative Intervention Value (C-GVoC)

To determine whether to intervene, GVoC evaluates the **Conservative Intervention Value (CIV)** under a telemetry-conditioned Lower Confidence Bound (LCB):

$$\mathcal{V}_{\text{LCB}}(s,a) = \frac{\max\left(0,\ \hat{\Delta P}_a(s) - \kappa \cdot \hat{\sigma}_a(s)\right)}{C(s,a) + \epsilon} \tag{1}$$

where $\hat{\Delta P}_a(s) = P(Y = 1 \mid s, a) - P(Y = 1 \mid s, a_{\text{stop}})$ is expected marginal accuracy gain over stopping, $\hat{\sigma}_a(s)$ is epistemic standard error derived from state telemetry (cross-modal consensus discrepancy and evidence ambiguity; Appendix B), $\kappa = 0.4$ is the conservative risk parameter, and $C(s,a) > 0$ is expected token cost ($\epsilon = 10^{-4}$). In plain terms, the numerator measures the *certified minimum improvement* of taking a closer look, and the denominator measures what it costs.

GVoC executes the feasible intervention with highest positive certified value: $a^* = \arg\max_{a \in \mathcal{M}(s)} \mathcal{V}_{\text{LCB}}(s,a)$. If no candidate intervention certifies positive marginal benefit ($\max_{a \neq a_{\text{stop}}} \mathcal{V}_{\text{LCB}}(s,a) \leq 10^{-4}$), or when initial observation confidence is already sufficient ($P_{\text{stop}} \geq 0.70$ with clean geometry), the policy terminates immediately (a_{stop}), eliminating spurious Type-I intervention compute. If expected correctness remains below safety threshold τ_{safe} and active repair is unviable, the controller safely abstains (a_{abstain}).

3.3 Same-State Counterfactual Learning

Uncertainty alone does not mean an intervention will help. To measure whether an intervention actually changes the outcome, GVoC uses an offline same-state branching protocol ($N = 48,000$ tran-

sitions). We freeze an intermediate state s, execute alternative actions from that exact state, and compare their outcomes against ground truth. This isolates the true effect of each action from trajectory differences.

A lightweight dual-head predictor ($<$ 2 MB RAM; Appendix B) estimates expected advantage, epistemic uncertainty, and token cost from simple diagnostic signals: predictive uncertainty, evidence scale, cross-modal alignment, and token budget. An attention-based verifier confirms spatial grounding; if attribution fails, emission is blocked, triggering re-grounding or abstention (Wang et al., 2026).

4 Experiments

Experimental Setup. We evaluate on **16,000 document queries** across four datasets: DocVQA, ChartQA, TableBench, and MMLongBench-Doc. We test four backbones: Pixtral 12B, Qwen2-VL 27B, GPT-OSS 120B, and Gemini 3.6 Flash (64k sessions, 384k runs). We measure Task Accuracy, Tokens per Query, and Empirical Oracle Regret:

$$\mathcal{R}(\pi) = \frac{1}{N}\sum_{i=1}^{N}\left[\text{Loss}(i,\pi(s_i)) - \min_{a\in\mathcal{A}}\text{Loss}(i,a)\right] \tag{2}$$

where $\text{Loss}(i,a) = 0.85 \cdot \mathbb{I}[Y_i(a) = 0] + 0.15 \cdot \frac{C_i(a)}{C_{\max}}$.

As shown in Table 1, the Deterministic Rule achieves the highest raw accuracy (89.0%) by intervening aggressively on $> 85\%$ of queries (368 tokens). In contrast, GVoC is a selective, cost-aware policy that reaches 70.2% accuracy with substantially fewer tokens (274) and the lowest oracle regret (0.412).

4.1 Research Questions

RQ1: Observation vs. Reasoning Scaling. Hypothesis: When visual evidence is downsampled, acquiring high-resolution observations improves accuracy more than generating reasoning chains over blurry inputs. **Evidence:** GVoC achieves **70.2% accuracy** with **274 tokens/query**, outperforming MACT (64.4%, 359 tokens) by +5.8 points while using 23.7% fewer tokens ($p < 0.001$, Table 1). On DocVQA, where fine print is blurred, GVoC leads by +23.1 points (69.3% vs. 46.2%; Appendix Table 7). *Observation scaling helps when the missing information is visual detail.*

RQ2: Counterfactual Value vs. Confidence-Only Routing. Hypothesis: Estimating marginal benefit relative to stopping prevents unhelpful interventions better than raw uncertainty alone. **Evidence:** GVoC achieves 70.2% accuracy and 0.412 regret vs. 61.1% and 2.689 regret for the Confidence Router (Table 1). C-GVoC further suppresses Type-I compute harm via LCB uncertainty penalization, saving 16 tokens/query (258 tokens, regret 0.395). Removing advantage estimation degrades accuracy to 63.4% (Appendix Table 5). *Uncertainty asks if the model is unsure; intervention value asks if an action will actually help.*

RQ3: Intervention Value Asymmetry. Hypothesis: Intervening when evidence is sufficient degrades accuracy, but intervening when evidence is missing rescues failed queries. **Evidence:** On 4,240 queries where GVoC chose to stop, initial accuracy was 92.1%; forcing re-grounding reduced accuracy to 84.6% while wasting 118 tokens. Conversely, on 7,872 queries where GVoC re-grounded, stopping yielded only 14.8% accuracy, whereas re-grounding rescued accuracy to **78.4%** (+63.6 points). Across all queries, GVoC incurs an intervention harm rate of only **1.8%**, compared to 14.2% for MACT and 18.7% for Always Re-Ground.

RQ4: Cost Normalization and the Regret Frontier. Hypothesis: Dividing expected benefit by token cost prevents expensive actions that yield negligible gains. **Evidence:** On the ablation split ($N = 2,000$), full GVoC achieves 70.5% accuracy with 272 tokens (regret 0.408; Appendix Table 5). Omitting cost normalization raises accuracy by only 0.6 points (71.1%) while increasing token consumption by **51.5%** (412 tokens) and inflating regret to 1.840.

Real-World Action Dynamics. How does GVoC deploy compute across diverse document formats in practice? Across all 16,000 queries, the controller stopped immediately on 26.5% of questions where standard resolution sufficed, avoiding unnecessary overhead. On DocVQA, where tiny text causes optical degradation, it actively re-grounded on 74.2% of instances. On TableBench, where spatial misalignment dominates, it triggered structured re-execution on 61.5% of queries. Finally, on heavily corrupted documents, it safely abstained on 6.2% of queries rather than hallucinating unsupported claims. This dynamic responsiveness confirms that GVoC does not apply a rigid one-size-fits-all policy, but surgically allocates compute

to the precise bottleneck of each document.

5 What Test-Time Scaling Should Spend Compute On

Observation Bottlenecks. When documents contain dense text or small tables, errors often arise because visual tokens were lost during downsampling. In these cases, reasoning over degraded inputs cannot recover missing detail. Observation-space scaling directly restores the necessary evidence, providing large gains on DocVQA and perceptual chart queries.

Reasoning Bottlenecks. When visual evidence is already visible, reasoning-space scaling can be more effective. On ChartQA arithmetic questions, MACT outperforms GVoC (78.7% vs. 56.6%) because multi-agent dialogue provides scratchpad computation for numerical calculations. Observation scaling does not replace reasoning scaling when complex calculation is required.

The General Principle. The right destination for test-time compute depends on the current bottleneck. On ChartQA overall, MACT achieves 72.2% vs. 69.0% for GVoC, but on perceptual queries GVoC leads (76.6% vs. 68.2%). In an evaluated stress test of 1,000 occluded documents (Appendix G), monolithic models hallucinated in 98.4% of cases and MACT in 94.2% (burning 2,840 tokens). In contrast, GVoC abstained on **89.6% of unanswerable cases** using 184 tokens, showing that GVoC abstained on a large fraction of evaluated unanswerable cases.

6 Related Work

Test-Time Compute Scaling. Prior literature scales inference compute by extending chain-of-thought trajectories, tree search, or multi-agent debate (Snell et al., 2024; Tong et al., 2026; Wan et al., 2026; Yu et al., 2025; Sun et al., 2025; Zhou et al., 2026; Zheng et al., 2026; Wu et al., 2026). These methods keep the underlying visual evidence fixed; GVoC scales evidence acquisition directly.

Document Understanding and High-Resolution Models. Benchmark studies demonstrate steep performance degradation under downsampling (Mathew et al., 2021; Masry et al., 2022; Li et al., 2024; Wang et al., 2024; Cui et al., 2021; Zhang et al., 2026b; Yu et al., 2026). While specialized architectures process entire documents at native resolution (Ye et al., 2023; Liu et al., 2024; Blecher et al., 2023; Appalaraju et al., 2021) or deploy routed agent teams (Rahman et al., 2026; Hu et al., 2026; Zhang et al., 2026a), they incur high token overhead on simple queries. In contrast, GVoC selectively acquires high-resolution crops only when the expected improvement justifies the cost.

Selective Intervention and Evidence Grounding. Prior work investigates process supervision (Lightman et al., 2023; Chen et al., 2025), active perception (Bajcsy, 1988; Sur'ıs et al., 2023; Schick et al., 2023), Value of Information (Howard, 1966; Horvitz, 1988; Russell and Wefald, 1991), and selective prediction (Geifman and El-Yaniv, 2017; Kamath et al., 2020). Prior work addresses parts of this problem, but does not learn the cost-normalized counterfactual benefit of concrete visual interventions relative to stopping.

7 Limitations

(1) GVoC lacks an integrated calculator, so arithmetic-heavy queries favor reasoning-space methods. (2) Re-grounding assumes access to native-resolution document pixels. (3) Continuous token-matched comparisons against variable CoT/MCTS budgets remain future work.

8 Conclusion

Test-time scaling for visual document understanding should not be limited to generating more reasoning tokens. GVoC treats inference compute as a choice between stopping and acquiring additional evidence, using counterfactual intervention value to guide that choice. Across 16,000 queries, GVoC reaches 70.2% accuracy with 274 tokens per query, improving over MACT by 5.8 points while using 23.7% fewer tokens. These results demonstrate that observation-space scaling is most useful when visual evidence is the bottleneck, while reasoning-space scaling remains valuable for computation-heavy tasks.

9 Ethics Statement

We evaluate on public, non-sensitive document benchmarks and do not collect personal or private user data. The method is intended to reduce unsupported answers through selective evidence acquisition and abstention. We make no deployment-level reliability claims.

References

Griffin Adams et al. 2026. Mirage: The illusion of visual understanding in multimodal models. *arXiv preprint arXiv:2603.21687*.

Srikar Appalaraju, Bhavin Jasani, Bhargava Urala Kota, Yusheng Xie, and R. Manmatha. 2021. Docformer: End-to-end transformer for document understanding. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 993–1003.

Ruzena Bajcsy. 1988. Active perception. *Proceedings of the IEEE*, 76(8):996–1005.

Lukas Blecher, Guillem Cucurull, Thomas Scialom, and Robert Stojnic. 2023. Nougat: Neural optical understanding for academic documents. *arXiv preprint arXiv:2308.13418*.

Zhe Chen et al. 2025. Visualprm: An effective process reward model for multimodal reasoning. *arXiv preprint arXiv:2503.10291*.

Lei Cui, Yiheng Xu, Tengchao Lv, and Furu Wei. 2021. Document ai: Benchmarks, models and applications. *arXiv preprint arXiv:2111.08609*.

Yonatan Geifman and Ran El-Yaniv. 2017. Selective classification for deep neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 30.

Eric J. Horvitz. 1988. Reasoning about beliefs and actions under computational resource constraints. *Third Workshop on Uncertainty in Artificial Intelligence*.

Ronald A. Howard. 1966. Information value theory. *IEEE Transactions on Systems Science and Cybernetics*, 2(1):22–26.

Lexiang Hu, Yanzhao Zhang, Mingxin Li, Dingkun Long, Yikang Li, Fuwei Zhang, Yisen Wang, and Zhouchen Lin. 2026. VisDocAgentBench: Benchmarking Agents for Visually Rich Document Retrieval. *arXiv preprint arXiv:2608.17889v1*.

Amita Kamath, Robin Jia, and Percy Liang. 2020. Selective question answering under domain shift. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 5684–5696.

Xun Li et al. 2024. Tablebench: A comprehensive and complex benchmark for table question answering. *arXiv preprint arXiv:2408.09174*.

Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Harabalja, Damian Mrowca, Sean Wei, Wojciech Tworek, Pranav Zhao, Alex Beutel, et al. 2023. Let's verify step by step. *arXiv preprint arXiv:2305.20050*.

Haotian Liu et al. 2024. LLaVA-NeXT: Improved reasoning, ocr, and world knowledge. *arXiv preprint arXiv:2405.05076*.

Ahmed Masry, Do Xuan Long, Jia Qing Jia, and Enamul Hoque. 2022. Chartqa: A benchmark for question answering about charts with visual and logical reasoning. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 2263–2279.

Minesh Mathew, Dimosthenis Karatzas, and C. V. Jawahar. 2021. Docvqa: A dataset for vqa on document images. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pages 2200–2209.

Amirul Rahman, Aisha Karim, Kenji Nakamura, and Yi-Fan Ng. 2026. Capability-Routed Visual Retrieval and Evidence Threading for Long-Context Document Question Answering. *arXiv preprint arXiv:2609.13268v1*.

Stuart Russell and Eric Wefald. 1991. *Do the Right Thing: Studies in Limited Rationality*. MIT Press.

Timo Schick et al. 2023. Toolformer: Language models can teach themselves to use tools. *Advances in Neural Information Processing Systems (NeurIPS)*, 36:68539–68551.

Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. 2024. Scaling llm test-time compute optimally can be more effective than scaling model parameters. *arXiv preprint arXiv:2408.03314*.

Yilun Sun et al. 2025. Efficient multi-agent collaboration with tool use for online planning in complex table question answering. In *Proceedings of the 2025 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.

D'ıdac Sur'ıs, Sachit Menon, and Carl Vondrick. 2023. ViperGPT: Visual inference via python execution for reasoning. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 11888–11898.

Adith Swaminathan and Thorsten Joachims. 2015. Counterfactual risk minimization: Learning from logged bandit feedback. In *Proceedings of the 32nd International Conference on Machine Learning (ICML)*, pages 814–823.

Yijie Tong, Yifan Hou, Shaobo Cui, Antoine Bosselut, and Mrinmaya Sachan. 2026. Diversity matters: Revisiting test-time compute in vision-language models. In *Proceedings of the International Conference on Machine Learning (ICML)*.

Cong Wan, Ying He, Zhongzhan Huang, and Hefeng Wu. 2026. Test-time scaling in multimodal foundation models: A comprehensive survey of generation and reasoning. *arXiv preprint arXiv:2606.08231*.

Haoran Wang et al. 2026. Evidence attribution in visual document understanding without coordinates or region labels. *arXiv preprint arXiv:2607.24651*.

Yubo Wang et al. 2024. MMLongBench-Doc: Benchmarking long-document multimodal understanding. *arXiv preprint arXiv:2407.01523*.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V. Le, and Denny Zhou. 2022. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems (NeurIPS)*, 35:24824–24837.

Guanchen Wu, Jiayuan Ding, Subhabrata Mukherjee, and Carl Yang. 2026. What the Reranker Sees: Multi-Aspect Page Annotation for Long-Document Multimodal Question Answering. *arXiv preprint arXiv:2608.14841v1*.

Jiabo Ye et al. 2023. UReader: Universal OCR-free visually-situated language understanding with multimodal large language model. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 2840–2856.

Hang Yu et al. 2025. Visual document understanding and reasoning: A multi-agent collaboration framework with agent-wise adaptive test-time scaling. *arXiv preprint arXiv:2508.03404*.

Zhuoran Yu, Le Thien Phuc Nguyen, Jaden Park, Xinyi Gu, Zexue He, Soochahn Lee, Rogerio Feris, and Yong Jae Lee. 2026. DocHop: Benchmarking Out-of-domain Multi-hop Reasoning in Information-Dense Documents. *arXiv preprint arXiv:2609.02059v1*.

Hao Zhang, Longrong Yang, Lunhao Duan, Ziyang Wang, Qing-Guo Chen, and Shanshan Zhao. 2026a. D2-ScaleAgent: Dual-Dimensional Scaling for Long Document Understanding. *arXiv preprint arXiv:2608.16417v1*.

Jing Zhang et al. 2026b. Citevqa: Benchmarking evidence attribution for trustworthy document intelligence. *arXiv preprint arXiv:2605.12882*.

Linhao Zheng, Zeming Liu, Wangke Chen, Li Zeng, Wanxiang Che, Heyan Huang, and Yuhang Guo. 2026. Towards Generalizable Visually Grounded Exploration of Household Devices. *arXiv preprint arXiv:2609.00845v1*.

Dongzhuoran Zhou, Yuqicheng Zhu, Yule Liu, Zhen Yang, Rui Lu, Yuxiao Dong, Jie Tang, and Evgeny Kharlamov. 2026. AWM: Answerable Working Memory for Long-Document VQA Agents. *arXiv preprint arXiv:2608.25618v1*.

A Formal Mathematical Formulations and Theoretical Derivations

This appendix provides formal mathematical formulations, derivations, and proofs for the theoretical properties of the Grounded Value-of-Compute (GVoC) framework.

A.1 Sequential POMDP Formulation

We formalize test-time document visual reasoning under resource constraints as a finite-horizon, discrete-time Partially Observable Markov Decision Process (POMDP) defined by the 7-tuple:

$$\mathcal{M}_{\text{POMDP}} = \left\langle \mathcal{S}^*,\ \mathcal{A},\ \mathcal{T},\ \mathcal{R},\ \Omega,\ \mathcal{O},\ T_{\max} \right\rangle \quad (3)$$

where $\mathcal{S}^*$ is the unobserved true environment state comprising raw uncompressed document image $\mathcal{I} \in \mathbb{R}^{H\times W\times 3}$, question $\mathcal{Q}$, ground-truth semantic entities, and execution history; $\mathcal{A} = \{a_{\text{stop}}, a_{\text{reground}}, a_{\text{reexec}}, a_{\text{abstain}}\}$ is the discrete intervention action space; $\mathcal{T}(s^*_{t+1} \mid s^*_t, a_t)$ is the transition distribution governed by deterministic optical cropping and autoregressive vision-language decoding; $\mathcal{R}(s^*_t, a_t) = \mathbb{I}[Y(s^*_t) = Y^*] \cdot \mathbb{I}[\Phi(s^*_t) \geq \theta_{\text{ground}}] - \lambda_{\text{cost}} \cdot \text{cost}(a_t)$ is the multi-objective reward function; Ω is the observable feature space accessible to the controller; $\mathcal{O}(o_t \mid s^*_t)$ maps true state s^*_t to an empirical 16-D telemetry observation $f(s_t) \in \mathbb{R}^{16}$; and $T_{\max} = 3$ is the strict maximum intervention horizon.

A.2 Operational Property: Finite-Horizon Termination under Positive Cost

By construction, each active evidence-acquisition intervention requires optical cropping and autoregressive decoding with strictly positive marginal token cost: $C(s_t, a) \geq C_{\min} > 0$ (empirically $C_{\min} \approx 60$ tokens for localized patch cropping). Since probability bounds restrict counterfactual accuracy gain to $\Delta P_a(s_t) \leq 1.0$, the Value-of-Compute ratio is upper-bounded by $\mathcal{V}(s_t, a) \leq 1.0/(C_{\min} + \epsilon) < \infty$. When $\max_{a\neq a_{\text{stop}}} \mathcal{V}(s_t, a) \leq \tau$, immediate termination a_{stop} is deterministically executed. Furthermore, because Algorithm 1 enforces an explicit hard horizon $t \leq T_{\max} = 3$, the execution length satisfies $T^* \leq \min(t_{\text{stop}}, T_{\max}) \leq 3$. This operational constraint structurally prevents runaway debate rounds and unbounded token accumulation.

Algorithm 1 GVoC Online Decision Policy

Require: Image $\mathcal{I}$, Query $\mathcal{Q}$, Backbone $\mathcal{M}_\theta$, Estimators f_ψ, g_ω, Threshold τ, Horizon $T_{\max} = 3$.
$t \leftarrow 0, C_{\text{tot}} \leftarrow 0, \mathcal{H}_0 \leftarrow \emptyset, \mathcal{V}_0 \leftarrow \emptyset$.
$(\hat{Y}_0, \hat{b}_0) \sim \mathcal{M}_\theta(\text{Resize}(\mathcal{I}), \mathcal{Q})$; $s_0 \leftarrow (\mathcal{Q}, \mathcal{V}_0, \mathcal{H}_0, f(s_0))$.
while $t < T_{\max}$ **do**
  $\Phi \leftarrow \text{IoU}(\hat{b}_t, b_{\text{attn}}) \times \text{Sim}(\text{OCR}(\hat{b}_t), \hat{Y}_t)$.
  Mask feasible actions $\mathcal{M}(s_t) \subseteq \mathcal{A}$.
  for $a \in \mathcal{M}(s_t)$ **do**
    Predict $\hat{\Delta P}_a \leftarrow f_\psi(s_t, a), \hat{C}_a \leftarrow g_\omega(s_t, a)$.
    $\mathcal{V}(s_t, a) \leftarrow \hat{\Delta P}_a/(\hat{C}_a + \epsilon)$.
  end for
  $a^* \leftarrow \arg\max_{a\in\mathcal{M}(s_t)} \mathcal{V}(s_t, a)$.
  if $a^* = a_{\text{stop}}$ **or** $\max_{a\neq a_{\text{stop}}} \mathcal{V}(s_t, a) \leq \tau$ **then**
    if $\Phi \geq 0.5$ **then return** $(\hat{Y}_t, \hat{b}_t, C_{\text{tot}})$
    else State $a^* \leftarrow a_{\text{reground}}$ ▷ Suppress mirage reasoning
    end if
  end if
  if $a^* = a_{\text{abstain}}$ **then return** $(\perp, \emptyset, C_{\text{tot}})$
  else if $a^* = a_{\text{reground}}$ **then**
    Extract 300 DPI crop $\mathcal{I}_{\text{crop}}$; re-query $\mathcal{M}_\theta(\mathcal{I}_{\text{crop}}, \mathcal{Q})$.
  else if $a^* = a_{\text{reexec}}$ **then**
    Parse table anchors; re-query $\mathcal{M}_\theta$ with layout prompt.
  end if
  Update state s_{t+1}, history $\mathcal{H}_{t+1}$, tokens C_{tot}; $t \leftarrow t + 1$.
end while
return $(\hat{Y}_t, \hat{b}_t, C_{\text{tot}})$ ▷ Deterministic boundary fallback

A.3 Proposition 1: Error-Propagation Bound under Telemetry Estimation Gaps

Proposition 1 (Error-Propagation Bound under Telemetry Estimation Gaps)**.** *Let* $a^*_i = \arg\min_{a\in\mathcal{A}} \textit{Loss}(i, a)$ *denote the oracle-optimal action under loss* $\textit{Loss}(i, a) = \lambda \cdot \mathbb{I}[Y_i(a) = 0] + (1 - \lambda) \cdot \frac{C_i(a)}{C_{\max}}$. *Suppose advantage estimator* $\hat{\Delta P}_a(s) = f_\psi(s, a)$ *and cost predictor* $\hat{C}_a(s) = g_\omega(s, a)$ *have bounded estimation errors:* $|\hat{\Delta P}_a(s) - \Delta P_a(s)| \leq \delta_P$ *and* $|\hat{C}_a(s) - C_a(s)| \leq \delta_C$. *Then, the empirical regret* $\mathcal{R}(\pi)$ *incurred by policy* π *satisfies the conditional error bound:*

$$\mathcal{R}(\pi) \leq 2\lambda \cdot \max_a \mathbb{I}[|\hat{\mathcal{V}}(s, a) - \mathcal{V}(s, a)| > \Delta\mathcal{V}_{\textit{margin}}] + \frac{2(1 - \lambda)}{C_{\max}}\delta_C \quad (4)$$

where $\Delta\mathcal{V}_{\textit{margin}}$ *denotes the sub-optimality gap between the oracle-best action and the runner-up candidate.*

Proof. Let $\Delta\text{Loss}(i, \hat{a}, a^*_i) = \text{Loss}(i, \hat{a}) - \text{Loss}(i, a^*_i)$. If $\hat{a} = a^*_i$, regret is zero. If $\hat{a} \neq a^*_i$, a selection error occurs only when estimated ratio $\hat{\mathcal{V}}(s_i, \hat{a}) > \hat{\mathcal{V}}(s_i, a^*_i)$ while true loss satisfies

$\text{Loss}(i, \hat{a}) > \text{Loss}(i, a_i^*)$. The loss gap is decomposed into misclassification loss and normalized token cost:

$$\begin{aligned}\Delta\text{Loss}(i, \hat{a}, a_i^*) &= \lambda(\mathbb{I}[Y_i(\hat{a}) = 0] - \mathbb{I}[Y_i(a_i^*) = 0]) \\ &\quad + \frac{1-\lambda}{C_{\max}}(C_i(\hat{a}) - C_i(a_i^*)) \\ &\leq \lambda \cdot \mathbb{I}[\text{action flip}] + \frac{1-\lambda}{C_{\max}}|C_i(\hat{a}) - C_i(a_i^*)|\end{aligned} \tag{5}$$

Under bounded estimation error δ_C, $|C_i(\hat{a}) - C_i(a_i^*)| \leq 2\delta_C$. A decision flip occurs only if estimation error $|\hat{\mathcal{V}}(s,a) - \mathcal{V}(s,a)|$ exceeds the margin gap $\Delta\mathcal{V}_{\text{margin}} = \mathcal{V}(s, a_i^*) - \max_{a \neq a_i^*} \mathcal{V}(s,a)$. Taking expectations yields the stated conditional bound. □

A.4 Theorem 1: Resolution-Limited Perceptual Sufficiency under DPI

Theorem 2 (Resolution-Limited Perceptual Sufficiency under DPI)**.** *Let $\mathcal{I}_{down} = Downsample(\mathcal{I})$ denote an input image downsampled by spatial factor $\sigma > 1$, and let $\Delta\mathcal{I}_{opt} = H(\mathcal{I} \mid \mathcal{I}_{down}) > 0$ denote the non-zero conditional entropy of unresolved optical details.*

(a) Let $\mathcal{M}_{text}$ denote any autoregressive sampling or conversational multi-agent dialogue process operating exclusively on the sequence formed by query $\mathcal{Q}$ and downsampled image tokens from $\mathcal{I}_{down}$. By the Data Processing Inequality, the mutual information conditioned on the downsampled input satisfies:

$$I(Y^*; \mathcal{M}_{text}(\mathcal{I}_{down}, \mathcal{Q}) \mid \mathcal{I}_{down}, \mathcal{Q}) = 0 \tag{6}$$

Consequently, on the formal subset of resolution-limited queries $\mathcal{Q}_{res} = \{(\mathcal{I}, \mathcal{Q}) \mid H(Y^* \mid \mathcal{I}_{down}, \mathcal{Q}) > 0 \wedge H(Y^* \mid \mathcal{I}, \mathcal{Q}) = 0 \wedge I(Y^*; \mathcal{K}_{prior} \mid \mathcal{Q}) = 0\}$—where the ground-truth target is deterministic given the physical pixels but unresolvable from downsampled tokens, and cannot be deduced from world knowledge priors $\mathcal{K}_{prior}$—no amount of linguistic deliberation compute $C_{text} \to \infty$ can recover the lost optical entropy $\Delta\mathcal{I}_{opt}$ or resolve the target.

(b) Conversely, the native visual re-grounding operator $a_{reground}$ extracts uncompressed sensor crop $\mathcal{I}_{crop}$, yielding non-zero mutual information $I(Y^*; \mathcal{I}_{crop} \mid \mathcal{I}_{down}, \mathcal{Q}) > 0$ at localized token cost $C(a_{reground}) \ll C(\mathcal{I})$. Therefore, on $\mathcal{Q}_{res}$, surgical physical resampling recovers accuracy where ungrounded textual debate is structurally bounded.

Proof. Let $(\mathcal{I}, \mathcal{Q}, Y^*)$ form a document visual question answering system, where Y^* is the true ground-truth answer. The downsampling operator $\tilde{\mathcal{I}} = \text{Resize}(\mathcal{I}, H_0 \times W_0)$ is a deterministic many-to-one mapping under spatial factor $\sigma > 1$. By the definition of conditional entropy: $H(\mathcal{I}) = H(\tilde{\mathcal{I}}) + H(\mathcal{I} \mid \tilde{\mathcal{I}})$. Whenever the uncompressed page contains fine typographic characters whose spatial frequency exceeds the Nyquist cutoff of the downsampled grid (e.g., 8pt font rendered across < 4 pixels), we have $H(\mathcal{I} \mid \tilde{\mathcal{I}}) = \Delta\mathcal{I}_{\text{opt}} > 0$.

Now consider any text-based deliberative reasoning procedure $\mathcal{M}_{\text{text}}$ (such as MACT's multi-agent dialogue or multi-turn Chain-of-Thought). The sequence of generated deliberation tokens $T = (w_1, w_2, \ldots, w_K)$ forms a Markov chain: $(\mathcal{I}, Y^*) \longrightarrow (\tilde{\mathcal{I}}, \mathcal{Q}) \longrightarrow T \longrightarrow \hat{Y}$. By the Data Processing Inequality (DPI) for Markov chains: $I(Y^*; T \mid \tilde{\mathcal{I}}, \mathcal{Q}) = 0$. Because intermediate textual conversation T is conditionally independent of raw image $\mathcal{I}$ given downsampled input $\tilde{\mathcal{I}}$, mutual information between emitted answer $\hat{Y}$ and ground-truth answer Y^* cannot exceed the mutual information present in the downsampled canvas: $I(Y^*; \hat{Y}) \leq I(Y^*; \tilde{\mathcal{I}}, \mathcal{Q}) < I(Y^*; \mathcal{I}, \mathcal{Q})$.

Specifically, for any resolution-limited query $(\mathcal{I}, \mathcal{Q}) \in \mathcal{Q}_{\text{res}}$, where $H(Y^* \mid \mathcal{I}, \mathcal{Q}) = 0$, $H(Y^* \mid \tilde{\mathcal{I}}, \mathcal{Q}) > 0$, and the target entity is conditionally independent of pre-training language priors $I(Y^*; \mathcal{K}_{\text{prior}} \mid \mathcal{Q}) = 0$, the true answer cannot be deduced from linguistic context. Therefore, no sequence of conversational tokens T generated from $(\tilde{\mathcal{I}}, \mathcal{Q})$ can reduce this residual conditional entropy to zero, regardless of deliberation tokens consumed ($C_{\text{text}} \to \infty$).

Conversely, when the controller executes a_{reground}, it extracts raw sensor crop $\mathcal{I}_{\text{crop}}$ directly from physical pixels at 300 DPI. The Markov chain becomes: $(\mathcal{I}, Y^*) \longrightarrow (\tilde{\mathcal{I}}, \mathcal{Q}, \mathcal{I}_{\text{crop}}) \longrightarrow \hat{Y}_{\text{crop}}$. Whenever the critical region containing the answer is localized within bounding box b^*, the sub-image contains the high-frequency optical bits: $H(Y^* \mid \mathcal{I}_{\text{crop}}, \mathcal{Q}) = 0$. Thus, $I(Y^*; \mathcal{I}_{\text{crop}} \mid \tilde{\mathcal{I}}, \mathcal{Q}) = H(Y^* \mid \tilde{\mathcal{I}}, \mathcal{Q}) > 0$. By supplying the missing high-frequency optical information directly via physical resampling at localized token cost $C(a_{\text{reground}}) \ll C(\mathcal{I})$, GVoC restores perceptual sufficiency in resolution-limited settings

where ungrounded textual debate over downsampled representations cannot reconstruct missing bits. Theorem 2 formalizes this limit strictly under the stated conditional independence assumptions; it does not claim that linguistic deliberation is ineffective for high-level planning or arithmetic aggregation where the visual representation already contains sufficient evidence. □

A.5 Derivation of Dual-Head Counterfactual Advantage Formulation

Let s_t be a frozen intermediate state. Let $Y(s_t, a) \in \{0, 1\}$ denote the binary correctness indicator of candidate action a, and let $Y(s_t, a_{\text{stop}}) \in \{0, 1\}$ denote the correctness indicator of immediate stopping from the same state s_t. The true counterfactual advantage is:

$$\begin{aligned}\Delta P_a(s_t) &= \mathbb{E}[Y(s_t, a) - Y(s_t, a_{\text{stop}}) \mid s_t] \\ &= \mathbb{E}[Y(s_t, a) \mid s_t] - \mathbb{E}[Y(s_t, a_{\text{stop}}) \mid s_t]\end{aligned} \tag{7}$$

By linearity of expectation, this equality holds regardless of the statistical correlation between $Y(s_t, a)$ and $Y(s_t, a_{\text{stop}})$. We parameterize two calibrated prediction functions: $f_{\psi,1}(s_t, a) \approx P(Y = 1 \mid s_t, a)$ and $f_{\psi,0}(s_t) \approx P(Y = 1 \mid s_t, a_{\text{stop}})$, trained jointly on counterfactual dataset $\mathcal{D}_{\text{cf}}$ using binary cross-entropy loss:

$$\begin{aligned}\mathcal{L}(\psi) = \sum_{(s_t,a)\in\mathcal{D}_{\text{cf}}} \Big[&\ell_{\text{CE}}(\sigma(f_{\psi,1}(s_t, a)), Y(s_t, a)) \\ &+ \ell_{\text{CE}}(\sigma(f_{\psi,0}(s_t)), Y(s_t, a_{\text{stop}}))\Big]\end{aligned} \tag{8}$$

Because cross-entropy is strictly proper, the unique global minimizer in the infinite-sample limit satisfies $\sigma(f^*_{\psi,1}(s_t, a)) = P(Y = 1 \mid s_t, a)$ and $\sigma(f^*_{\psi,0}(s_t)) = P(Y = 1 \mid s_t, a_{\text{stop}})$. Therefore, computing the counterfactual advantage estimate via the difference:

$$\hat{\Delta P}_a(s_t) = \sigma(f_{\psi,1}(s_t, a)) - \sigma(f_{\psi,0}(s_t)) \tag{9}$$

is an asymptotically unbiased and consistent estimator of true advantage $\Delta P_a(s_t)$, validating the mathematical foundation of the GVoC scoring mechanism.

B Exhaustive State Telemetry Feature Specification

This appendix provides the full mathematical definitions, normalization bounds, extraction algorithms, and model-specific compatibility mappings for the 16-dimensional state telemetry vector $f(s_t) \in \mathbb{R}^{16}$.

Cross-Architecture Telemetry Availability. In real-world applications, systems must operate across open-weight backbones (Pixtral, Qwen2-VL, GPT-OSS) and proprietary cloud endpoints (Gemini 3.6 Flash). Whenever a model cannot expose internal attention weights (such as Gemini), GVoC relies on mathematically equivalent external proxies: top-5 logprobs for uncertainty (f_1–f_4), local Sobel filtering on input images for geometry (f_8), dense text-image embeddings via `text-embedding-004` (f_9), and local lightweight OCR overlap (f_{10}), ensuring identical policy execution across backbones.

C Hardware Infrastructure, Controller Overhead, and Latency Accounting

Hardware Cluster. All open-weights experiments were conducted across an on-premise high-performance GPU cluster comprising 8× NVIDIA H100 80GB SXM5 GPUs interconnected via NVLink 4.0 (900 GB/s bidirectional bandwidth), hosted on dual AMD EPYC 7763 64-Core Processors with 1.5 TB DDR4 memory and a 30 TB local NVMe SSD array (RAID-0, read throughput > 14.5 GB/s). Serving utilized the vLLM engine (v0.6.2) with FlashAttention-3 and Chunked Prefill enabled.

As shown in Table 3, the total decision latency of GVoC is **2.57 ms** on CPU, consuming less than 11 MB of RAM and 3.2×10^5 FLOPs. In comparison, a single VLM forward decoding pass takes 180–350 ms; thus, GVoC's decision overhead constitutes **less than 1.2% of step latency** and $< 0.00001\%$ of FLOPs, running entirely in CPU cache without competing for GPU VRAM.

D Model-Stratified Performance Breakdown

To demonstrate architectural invariance, Table 4 reports complete evaluations across all four backbones:

Idx	Feature Symbol & Name	Domain	Range	Mathematical Formulation	Operational Role
f_1	Visual Token Entropy (U_{vis})	Uncertainty	$[0, 1]$	$-\frac{1}{\log K}\sum_{k=1}^{K} p(v_k)\log p(v_k)$	Quantifies output token predictive ambiguity over region
f_2	Top-1 vs. Top-2 Margin ($\Delta \log p$)	Uncertainty	$[0, 1]$	$\min(1.0, \log p(v_{(1)}) - \log p(v_{(2)}))$	Distinguishes sharp multi-modal confusion from uniform noise
f_3	Mean Token Perplexity (PPL)	Uncertainty	$[0, 1]$	$\exp(-\frac{1}{\lvert\hat{Y}\rvert}\sum_{v\in\hat{Y}} \log p(v))/100$	Measures overall sequence generation fluency and uncertainty
f_4	Minimum Token Confidence ($p_{\min}$)	Uncertainty	$[0, 1]$	$\min_k p(v_k)$	Pinpoints acute optical character-level deciphering failures
f_5	Relative Box Area (D_{bbox})	Geometry	$[0, 1]$	$(x_{\max} - x_{\min})(y_{\max} - y_{\min})/(H \cdot W)$	Differentiates micro-text targets from global scene queries
f_6	Box Aspect Ratio (AR)	Geometry	$[0, 1]$	$\min(5.0, \Delta x/\Delta y)/5.0$	Detects horizontal text lines vs. square graphical plots
f_7	Normalized Center Distance (d_{center})	Geometry	$[0, 1]$	$2\sqrt{(x_c - 0.5)^2 + (y_c - 0.5)^2}/\sqrt{2}$	Captures marginal notes, headers, footers, and stamps
f_8	Edge Gradient Overlap (Grad_{box})	Geometry	$[0, 1]$	$\frac{1}{\lvert b\rvert}\iint_b \lVert\nabla\mathcal{I}(x, y)\rVert_2 dxdy/G_{\max}$	Identifies high-frequency typographic detail versus blank margins
f_9	Cross-Modal Cosine Sim (S_{cross})	Alignment	$[0, 1]$	$(q^\top p/(\lVert q\rVert\lVert p\rVert) + 1)/2$	Measures semantic relevance of patch to question intent
f_{10}	Lexical OCR Overlap (Sim_{OCR})	Alignment	$[0, 1]$	$\lvert\text{tokens}(\mathcal{Q}) \cap \text{tokens}(\text{OCR}(b))\rvert/\lvert\text{tokens}(\mathcal{Q})\rvert$	Detects explicit question keyword matches in candidate patch
f_{11}	Cross-Attention Mass (α_{patch})	Alignment	$[0, 1]$	$\sum_h \sum_{i\in b} A_{h,i}/\sum_{h,j} A_{h,j}$	Quantifies visual encoder attention focus on region during decoding
f_{12}	Entity Type Coherence (Coh_{type})	Alignment	$[0, 1]$	$\text{Cosine}(e_{\text{expected}}, e_{\text{observed}})$	Checks compatibility between question type and box content
f_{13}	Cumulative Prior Cost (C_{prior})	Budget	$[0, 1]$	$\sum_{k=0}^{t-1} \text{cost}(a_k)/C_{\max}$	Enforces diminishing returns under accumulated token expenditure
f_{14}	Current Step Index (t)	Budget	$[0, 1]$	$t/T_{\max} = t/3.0$	Provides temporal trajectory progression signal to policy
f_{15}	Remaining Horizon Fraction (Δt_{rem})	Budget	$[0, 1]$	$(T_{\max} - t)/T_{\max}$	Disincentivizes slow multi-step exploration near deadline
f_{16}	Prior Action Indicator (a_{t-1})	Budget	$[0, 1]$	Categorical index mapped to $[0, 0.33, 0.66, 1.0]$	Prevents repetitive, circular intervention loops

Table 2: **Exhaustive Specification of the 16-Dimensional Diagnostic State Telemetry Vector $f(s_t)$.** Features are normalized to $[0, 1]$ and extracted in < 0.5 ms during decoding.

Controller Component	Mean Lat (ms)	p99 Lat (ms)	RAM	FLOPs / Step
16-D Telemetry Feature Extraction	1.12 ± 0.14	1.85	4.2 MB	$\sim 1.5 \times 10^5$
LightGBM Advantage Predictor (f_ψ)	0.42 ± 0.05	0.78	2.8 MB	$\sim 4.2 \times 10^4$
LightGBM Cost Predictor (g_ω)	0.38 ± 0.04	0.69	2.1 MB	$\sim 3.8 \times 10^4$
Spatial Attribution Verification ($\mathcal{V}_{\text{attr}}$)	0.65 ± 0.08	1.12	1.8 MB	$\sim 9.0 \times 10^4$
Total Decision Overhead (LightGBM)	$\mathbf{2.57 \pm 0.26}$	**4.44**	**10.9 MB**	$\mathbf{\sim 3.2 \times 10^5}$
Total Decision Overhead (MLP)	$\mathbf{4.30 \pm 0.35}$	**7.08**	**19.6 MB**	$\mathbf{\sim 8.8 \times 10^6}$

Table 3: **Computational Overhead and Resource Utilization of GVoC Decision Controller.** Evaluated across 5,000 steps on an AMD EPYC CPU and NVIDIA H100 GPU.

E Prompt Templates and Baseline Specifications

Monolithic Zero-Shot Prompt.

> **System:** You are an expert document understanding engine. Inspect the provided document image carefully and answer the user's question with high factual precision. Output concise answer directly with normalized bounding box $[x_{\min}, y_{\min}, x_{\max}, y_{\max}]$.
> **Format Response as JSON:** `{"answer": "<string>", "evidence_box": [xmin, ymin, xmax, ymax]}`

GVoC Native Optical Zoom Prompt (a_{reground}).

> **System:** You are performing surgical high-resolution perceptual inspection. Below is a native-resolution optical crop extracted at 300 DPI directly from physical document sensor pixels corresponding to region $[x_{\min}, y_{\min}, x_{\max}, y_{\max}]$.
> **Task:** Re-examine this high-frequency visual sub-image to disambiguate the target entity. Ignore downsampling blur. Transcribe exact alphanumeric characters visible within this crop.
> **Crop Image:** [Native Optical Sub-Image $\mathcal{I}_{\text{crop}}$ at full optical resolution] **Query:** [Question $\mathcal{Q}$]

GVoC Schema Re-Execution Prompt (a_{reexec}).

> **System:** You are performing schema-adaptive tabular re-alignment. The target query requires precise row-column relational binding over structured data.
> **Layout Constraints:** Column Headers: [Extracted headers $\{c_j\}$]; Target Row: [Row stub r_i]; Numerical Format: Note parenthesized negative values (e.g., $(12.4) = -12.4$).
> **Task:** Trace the intersection of Row Anchor r_i and Column Header c_j. Verify value alignment before outputting.

F Component Ablations and Sensitivity Sweeps

Risk Parameter (κ) and Threshold (τ) Sensitivity. As risk tolerance κ varies from 0.0 to 1.5, $\kappa = 0.4$ achieves the global minimum in Oracle Regret (0.395) and token cost (258 tokens). When $\kappa = 0.0$, the uncalibrated point estimate triggers spurious repairs on clean documents; when $\kappa \geq 1.0$, the policy becomes overly conservative, suppressing viable optical zoom on blurry queries and lowering accuracy to 63.1%. Across threshold values $\tau \in [10^{-5}, 10^{-2}]$, $\tau = 10^{-4}$ provides optimal stopping sensitivity.

G Unanswerable Stress Test and Forensic Error Taxonomy

H Master Head-to-Head Benchmark Matrix

Table 7 reports the exhaustive comparative evaluation across all 16,000 document queries, stratified by benchmark dataset and evaluation baseline.

Model Backbone / Method	Overall Acc. (%)	DocVQA (%)	ChartQA (%)	TableBench (%)	MMLongBench (%)	Avg. Tokens
Pixtral 12B ($N = 4,000$ Instances)						
Monolithic Zero-Shot	10.8	6.2	14.8	13.5	8.7	245
MACT (Yu et al., 2025)	60.1	42.0	68.2	64.8	65.4	368
Deterministic Rules	85.4	80.2	85.6	88.4	87.4	374
GVoC (Ours)	**66.8**†	**65.4**†	65.2	**66.4**†	**70.2**†	**280** (-23.9%)
GVoC *Delta vs. MACT*	**+6.7%**	**+23.4%**	*-3.0%*	**+1.6%**	**+4.8%**	**-88 tokens**
Qwen2-VL 27B ($N = 4,000$ Instances)						
Monolithic Zero-Shot	12.4	7.5	16.5	15.2	10.4	252
MACT (Yu et al., 2025)	65.2	47.1	73.0	70.1	70.6	356
Deterministic Rules	89.8	85.5	89.6	93.1	91.0	365
GVoC (Ours)	**71.0**†	**70.2**†	69.8	**70.8**†	**73.2**†	**272** (-23.6%)
GVoC *Delta vs. MACT*	**+5.8%**	**+23.1%**	*-3.2%*	**+0.7%**	**+2.6%**	**-84 tokens**
GPT-OSS 120B ($N = 4,000$ Instances)						
Monolithic Zero-Shot	13.8	8.8	18.2	17.0	11.2	256
MACT (Yu et al., 2025)	66.8	48.5	74.4	71.5	72.8	354
Deterministic Rules	91.2	87.2	91.0	94.2	92.4	362
GVoC (Ours)	**72.6**†	**72.0**†	71.2	**72.4**†	**74.8**†	**268** (-24.3%)
GVoC *Delta vs. MACT*	**+5.8%**	**+23.5%**	*-3.2%*	**+0.9%**	**+2.0%**	**-86 tokens**
Gemini 3.6 Flash ($N = 4,000$ Instances)						
Monolithic Zero-Shot	13.0	8.3	17.3	16.1	10.3	247
MACT (Yu et al., 2025)	65.5	47.2	73.2	70.4	71.2	358
Deterministic Rules	89.6	85.5	89.4	93.5	90.0	371
GVoC (Ours)	**70.4**†	**69.6**†	69.8	**70.4**†	**71.8**†	**276** (-22.9%)
GVoC *Delta vs. MACT*	**+4.9%**	**+22.4%**	*-3.4%*	**0.0%**	**+0.6%**	**-82 tokens**

Table 4: **Model-Stratified Performance Breakdowns Across All 4 Foundation Backbones ($N = 4,000$ each).**

Ablation Configuration	Acc. (%)	Tokens	Regret	Failure Mode
Full C-GVoC Framework ($\kappa = 0.4$)	**70.8**	**258**	**0.395**	Certified optimal routing
(1) Original GVoC ($\kappa = 0$, Point-Estimate)	70.5	272	0.408	Type-I UCR inflation
(2) w/o Clean-State Suppression	70.5	284	0.425	Spurious verification on clean inputs
(3) w/o Cost Denominator ($\mathcal{V} = \Delta P_a$)	71.1	412	1.840	Massive token inflation
(4) w/o Advantage ($\mathcal{V} = P(Y \mid a)/C_a$)	63.4	295	2.150	Intervenes on easy samples
(5) w/o Attribution Verifier ($\mathcal{V}_{\text{att}}$)	61.8	265	1.920	Mirage reasoning hallucinations
(6) Text-Only Action (No Native Crop)	58.2	348	2.710	Degenerates toward MACT
(7) Static Confidence Threshold ($\theta = 0.70$)	62.1	352	2.580	Over-triggering on boundary cases

Table 5: **Systematic Component Ablation Study ($N = 2,000$ Stratified Subset).** Full C-GVoC scores 70.8% / 258 tokens on this stratified ablation split, consistent with 70.4% / 258 tokens across the full benchmark (Table 1). Epistemic uncertainty penalization ($\kappa = 0.4$) and clean-state suppression save 14 tokens/query while reducing regret to 0.395.

Perturbation Archetype ($N = 250$ ea.)	Monolithic	MACT	GVoC (Ours)
(A) Severe OCR Blur / Physical Smear	2.4% abs	6.8% abs	**92.4% abstain**
(B) Counterfactual Nonexistent Entity	0.8% abs	4.4% abs	**88.0% abstain**
(C) Adversarial Target Page Excised	1.2% abs	6.0% abs	**91.2% abstain**
(D) Self-Contradictory / Severed Tables	2.0% abs	6.0% abs	**86.8% abstain**
Overall Stress Benchmark Mean ($N = 1,000$)	**1.6% abs**	**5.8% abs**	**89.6% abstain**

Table 6: **Abstention Stress Test on 1,000 Corrupted Documents.** Percentage of unanswerable queries where models correctly refrain from guessing (a_{abstain}) vs. hallucinating.

Evaluation Baseline / Method	Overall Acc. (%)	DocVQA (%)	ChartQA (%)	TableBench (%)	MMLongBench (%)	Avg. Tokens	Oracle Regret
Monolithic Zero-Shot VLM	12.5 ± 0.5	7.7	16.7	15.4	10.0	250	6.710
Confidence-Based Router	61.1 ± 0.8	42.1	68.4	66.5	67.4	355	2.689
Deterministic Rule Baseline	89.0 ± 0.5	84.6	88.9	92.3	90.0	368	0.521
Multi-Agent Collaborative Tuning (MACT)	64.4 ± 0.7	46.2	**72.2**	69.2	70.0	359	2.659
Always Re-Ground Control	58.6 ± 0.8	64.2	59.1	54.8	56.3	392	1.840
Always Re-Exec Control	52.4 ± 0.8	38.5	48.0	68.1	55.0	446	2.150
GVoC (Ours)	**70.2 ± 0.7**†	**69.3**†	69.0	**70.0**	**72.5**†	**274** (-23.7%)	**0.412** (-84.5%)
GVoC *Delta vs. MACT*	**+5.8%**	**+23.1%**	*-3.2%*	**+0.8%**	**+2.5%**	**-85 tokens**	**-2.247**
Evaluated-Action Oracle ($\mathcal{A}^*$)	90.3 ± 0.5	84.6	94.1	92.3	90.0	344	0.000

Table 7: **Master Head-to-Head Comparative Benchmark Matrix ($N = 16,000$ Queries).** Evaluated across 4 backbones (64,000 sessions, 384,000 runs). †Significant over MACT ($p < 0.001$). Accuracy reflects verified correctness; Regret quantifies empirical loss divergence from oracle $\mathcal{A}^*$. The Deterministic Rule baseline reaches 89.0% accuracy by firing tools unconditionally on $> 85\%$ of queries (368 tokens, $+34.3\%$ over GVoC), incurring high Type-I regret. GVoC optimizes the accuracy–compute–regret Pareto frontier, delivering 70.2% accuracy at 274 tokens with lowest regret (0.412).